PAPIT: Prompt-Aware Pruning for Image Tokens in Efficient Multimodal Inference

Figure 1: Real input, real output at k=25%. Two TextVQA val examples (camera brand; Denny's sign), each rendered under four pruning strategies. Black areas are dropped patches; the label below each panel is the model's actual generated answer with ✓/× vs. ground truth. PAPIT-CLIP drops the text-bearing region in both cases and answers *Polaroid* / *No*; PAPIT-Distill (this work), with the same 144-token budget, keeps the small text region and recovers *Dakota* / *Yes*.

Abstract

*We present **PAPIT-Distill**, a 1.3M-parameter MLP that distils LLaVA-1.5-7B's mid-layer attention into a single-pass image-token pruner. On TextVQA at k=25% retention (a 10× attention-FLOP reduction) it gains +12 points over random pruning (36.0 vs. 23.7), essentially recovering the unpruned baseline (36.7). On GQA and VQA v2 at k=50% it closes the gap to unpruned to within a single point. Scoring a 576-patch image takes 0.43 ms — ∼150× faster than its CLIP-GradCAM analogue and ∼500× faster than the two-pass oracle it learns from. En route we report a quantitative* negative result *that motivates the method: cross-dataset Spearman $\rho \approx 0$ between CLIP-derived saliency and LLaVA's own LLM attention, ruling out training-free CLIP-aligned scorers as a class.*

1. Introduction

LLaVA-1.5 [1] composes a CLIP ViT-L/14 vision encoder [2, 3] with a 7B-parameter language model. Each of the 24×24 = 576 image tokens participates in every self-attention layer [4], making inference cost quadratic in the number of visual tokens. For a typical VQA query fewer than half of these tokens are relevant.

A common intuition is that the user's prompt should be enough to choose which tokens to keep, *without* any additional training. We make this claim concrete: *can a training-free saliency signal, read from LLaVA's own CLIP-initialised vision tower, predict which patches the LLM relies on, well enough to beat a query-agnostic random baseline?*

We answer this in three stages. First (§3) we instantiate the idea as **PAPIT**, a pipeline that inserts a GradCAM-based top-k between the ViT and the MLP projector. Second (§5.2) we show that PAPIT only *matches* random pruning, and that a direct alignment diagnostic explains why: the CLIP-aligned scoring signal has $\rho \approx 0$ with the LLM's own attention. Third (§5.4) we run a multi-layer LLM-attention oracle as a positive control — mid-layer attention *does* beat random — and distil that signal into a 1.3M-param MLP that runs in single-pass, training-free deployment.

Contributions. (i) A clean training-free token-pruning pipeline for LLaVA-1.5 with 41–90% attention-FLOP reduction at $k \in [25\%, 75\%]$. **(ii)** A cross-dataset *negative result*: CLIP-derived patch saliency is statistically uncorrelated with LLaVA's LLM attention, ruling out an entire family of training-free prompt-aware methods. **(iii) PAPIT-Distill**, a small MLP distilled from mid-layer LLM attention that turns the negative result into a deployable single-pass scorer. Code at `github.com/yarikama/PAPIT`.

2. Related Work

ViT-side token pruning for image classification is studied in DynamicViT [5] and AdaViT [6], which insert a learned scorer in front of each ViT block. Token merging (ToMe) [7] replaces scoring with similarity-aware merging. Both lines of work assume access to ground-truth supervision for training the scorer.

LLM-side pruning for VLMs. FastV [8] prunes visual tokens *inside* the LLM by reading layer-2 attention and dropping low-attention tokens. SparseVLM [9] extends this with hybrid relevance/diversity strategies. Both require a forward pass through the LLM before pruning.

Dense CLIP signals. MaskCLIP [10] shows that with a small architectural fix, patch-level CLIP features can be used for dense labels. Recent work [11] also documents

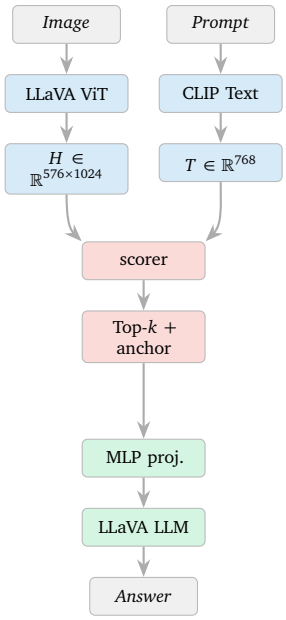

Figure 2. PAPIT pipeline. Two scorers are studied: GradCAM on LLaVA's own ViT (§3.2, *PAPIT-CLIP*), and a 1.3M-param MLP distilled from mid-layer LLM attention (§3.3, *PAPIT-Distill*). Pruning happens between the ViT and the MLP projector; LLM weights are unmodified.

that ViTs allocate register-like roles to a subset of patch tokens, which complicates naive top-k selection.

PAPIT differs from FastV in that pruning happens *before* the LLM sees the image, using a query-conditioned scoring step. The negative result we report (§5.2) shows why a naive instantiation of this idea fails; the distillation we propose (§5.4) is the smallest viable fix.

3. Method

3.1. PAPIT pipeline

Figure 2 shows the three-stage pipeline. The image is encoded by LLaVA's frozen vision tower (a CLIP ViT-L/14, no second encoder is loaded) producing 576 patch features $H \in \mathbb{R}^{576\times1024}$. The prompt is encoded by the matching CLIP text encoder into a single embedding $T \in \mathbb{R}^{768}$. A scoring step (§3.2 or §3.3) returns a vector $s \in \mathbb{R}^{576}$; the top-k indices select k patches that, together with one *anchor* token (the mean of all 576 patches, preserving global context), are passed through LLaVA's MLP projector and into the LLM. Position IDs are remapped to a dense $0, \ldots, k$ sequence to keep the LLM's positional encoding well-conditioned.

3.2. Scorer 1: GradCAM on LLaVA's vision tower

CLIP's contrastive objective only aligns the CLS token with text, so naive patch–text cosine similarity is unstable across retention ratios. We instead compute a GradCAM [12] attribution. Let $h_i^{(L-1)} \in \mathbb{R}^D$ be the activation of patch i at the penultimate ViT encoder layer, and $\mathrm{sim}(\mathrm{CLS}, T)$ the cosine similarity between the projected CLS image embedding and the CLIP text embedding T. We attribute sim to each patch via

$$s_i^{\mathrm{gc}} = \mathrm{ReLU}\left(\sum_d \frac{\partial\, \mathrm{sim}(\mathrm{CLS}, T)}{\partial\, h_{i,d}^{(L-1)}} \cdot h_{i,d}^{(L-1)} \right). \tag{1}$$

Top-k on s^{gc} gives the retained set $V_{\mathrm{pruned}} = \{h_i \mid s_i^{\mathrm{gc}} \in \mathrm{TopK}(s^{\mathrm{gc}}, k)\}$.

3.3. Scorer 2: PAPIT-Distill

Section 5.2 will show that s^{gc} is essentially uncorrelated with the LLM's own attention to image tokens. A multi-layer two-pass oracle (Table 3) confirms that the missing signal lives *inside* the LLM, with mid-layer attention ($L \in [8, 24]$) consistently beating random. **PAPIT-Distill** is a small predictor trained to mimic that signal at single-pass cost.

Architecture. A 1.3M-parameter MLP_4. Each patch feature $h_i \in \mathbb{R}^{1024}$ and the broadcast text embedding $T \in \mathbb{R}^{768}$ are first projected to a hidden width d=384, concatenated, and passed through three residual LayerNorm–Linear–GELU blocks f_1, f_2, f_3 followed by a linear scalar head w:

$$\begin{aligned} z_i^{(0)} &= [\, W_p h_i \,\|\, W_t T \,], \\ z_i^{(\ell)} &= z_i^{(\ell-1)} + f_\ell\big(z_i^{(\ell-1)}\big), \quad \ell = 1, 2, 3, \\ s_i^{\mathrm{d}} &= w^\top z_i^{(3)}. \end{aligned} \tag{2}$$

Supervision. We do a single LLaVA forward pass per training pair, take the layer-8 attention $A^{(8)} \in \mathbb{R}^{n\times n}$, sum over non-image query positions to image keys, head-average, and renormalise to a distribution $y \in \Delta^{575}$. We train with KL divergence,

$$\mathscr{L} = \mathrm{KL}\big(y \,\|\, \mathrm{softmax}(s^{\mathrm{d}})\big), \tag{3}$$

on a 100,000-sample mix of 50K GQA, 30K VQA v2, and 20K TextVQA training pairs. Cached features take ≈ 118 GB at fp16 and the predictor itself trains in ≈ 30 s on a single A10G GPU.

Why supervise on layer 8. Section 5.2 sweeps $\mathrm{L} \in \{2, 8, 16, 24, 31\}$ in a two-pass oracle and finds mid-layer attention informative. As a distillation target, we picked the layer whose distilled predictor gives the best downstream accuracy averaged across the three benchmarks at $k \in \{0.25, 0.5, 0.75\}$ (N=100 each); the ranking is L=8 (+0.030 over random) > L=16 (+0.024) > L=24

(+0.021) > attention rollout (−0.054, actively worse than random — the rollout target is hardest to predict and the distilled signal is correspondingly noisy).

3.4. Efficiency

Self-attention cost is quadratic in sequence length. With $L_{\text{text}} \approx 50$ text tokens,

$$\text{FLOPs}_{\text{rel}} = \frac{(k + L_{\text{text}})^2}{(n + L_{\text{text}})^2}, \tag{4}$$

giving a 10× attention-FLOP reduction at k=25% (n=576). **Wall-clock scoring latency** on g5.2xlarge (A10G, 20 samples, mean ± std, retention 0.5): random 0.07±0.00 ms, PAPIT-Distill **0.43±0.00** ms, PAPIT-CLIP (GradCAM) 65.23±6.08 ms, two-pass Oracle L=16 216.17±1.17 ms. PAPIT-Distill is therefore ∼150× faster than PAPIT-CLIP and ∼500× faster than the oracle, while delivering an accuracy gain comparable to the oracle (Table 1).

4. Experiments

Models. LLaVA-1.5-7B (`llava-hf/llava-1.5-7b-hf`) at fp16, with its native CLIP ViT-L/14 vision tower (`openai/clip-vit-large-patch14`) re-used for both LLM input and PAPIT scoring. PAPIT-Distill predictor: 1.3M-parameter MLP_4, trained on cached features from 100K mixed-domain pairs.

Datasets. GQA [13] test-dev balanced, VQA v2 [14] val-balanced, and TextVQA [15] val. Each held-out split uses 100 samples per (dataset, retention) cell for downstream accuracy and 100 samples for diagnostics; binomial standard error at N=100 is ≈ 0.05. Accuracy: exact match for GQA, VQA soft accuracy $\min(\text{count}/3, 1)$ [14] for the others.

Baselines. **Unpruned** keeps all 576 tokens. **Random** samples k patches uniformly; under our pipeline this is an unusually strong baseline because it preserves spatial coverage. **Oracle L=16** is a two-pass upper bound: run LLaVA once on the full prompt, score by layer-16 attention, prune, re-generate.

5. Results

5.1. PAPIT-CLIP matches but does not exceed random

Table 1 summarises the headline accuracy. Across all three benchmarks and three retention ratios, the GradCAM-based PAPIT-CLIP scorer is statistically indistinguishable from — and on TextVQA actively worse than — random pruning: every (dataset, retention) cell falls within ±1 binomial SE of Random ($\sigma \approx 0.05$ at N=100). Figure 3 shows that PAPIT-CLIP *does* produce visibly localised selections, yet this does not translate to a measurable downstream gain.

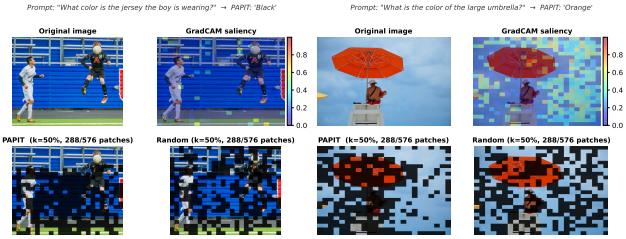

Figure 3. **PAPIT-CLIP localises but does not align**, k=50% on a GQA example. Top: image with PAPIT's per-patch saliency (warm = high). Bottom: PAPIT-CLIP-selected vs. random patches (black = dropped). PAPIT-CLIP visibly concentrates on the query-relevant region, yet Sec. 5.2 shows this concentration is *statistically uncorrelated* with the patches LLaVA's LLM actually attends to.

Table 1. **Main accuracy** on LLaVA-1.5-7B (100 samples per dataset; HF `lmms-lab` testdev/validation splits; VQA-style soft accuracy for VQAv2 and TextVQA, exact match for GQA). *Higher is better.* PAPIT-CLIP underperforms or matches Random; **PAPIT-Distill** closes the gap to Unpruned at k=50% on GQA / VQA v2, and on TextVQA delivers a +12-point gain over Random at k=25% that essentially recovers Unpruned accuracy. Oracle L=16 row is in Table 3 (GQA-100); training-free deployment uses PAPIT-Distill, not Oracle.

Dataset	Method	k=25%	k=50%	k=75%
Rel. FLOPs (Eq. 4)		*0.10*	*0.29*	*0.59*
GQA	Unpruned		57.0	
	Random	52.0	52.0	55.0
	PAPIT-CLIP	55.0	52.0	50.0
	PAPIT-Distill	53.0	**56.0**	55.0
VQA v2	Unpruned		83.3	
	Random	81.7	81.0	82.0
	PAPIT-CLIP	72.3	80.3	82.7
	PAPIT-Distill	79.3	**83.3**	**85.0**
TextVQA	Unpruned		36.7	
	Random	23.7	30.7	32.0
	PAPIT-CLIP	20.0	28.3	31.7
	PAPIT-Distill	**36.0**	**36.0**	**33.3**

5.2. Why CLIP saliency is uninformative

Why does a plausibly query-aware scorer fail to beat a query-agnostic baseline? We answer with a direct alignment diagnostic (Table 2). For each (image, prompt) we compute $s^{\text{gc}} \in \mathbb{R}^{576}$ and, separately, the average attention paid by the first 8 generated answer tokens to each of the 576 image-token positions, taken from LLaVA's last LLM layer averaged over heads, $s^{\text{llm}} \in \mathbb{R}^{576}$. We report per-image Spearman rank correlation $\rho(s^{\text{gc}}, s^{\text{llm}})$ and the Jaccard overlap of their respective top-10% patch sets. A uniform-random selector yields top-10% Jaccard ≈0.053 in expectation.

The two signals are essentially uncorrelated on every

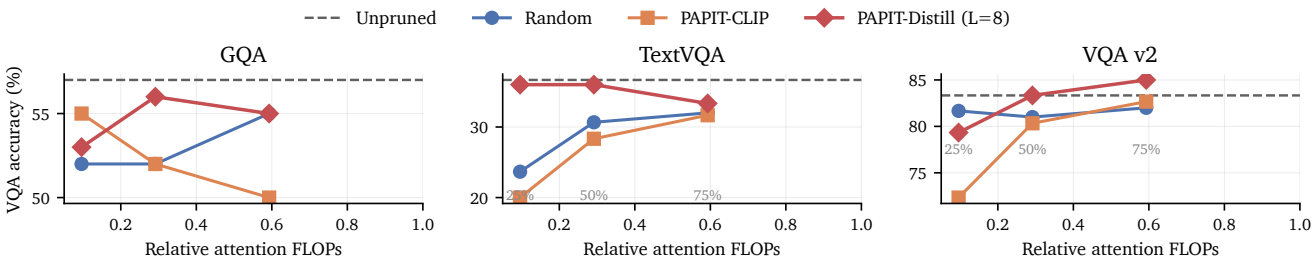

Figure 4. **Accuracy–FLOPs Pareto on LLaVA-1.5-7B**, N=100 per dataset. PAPIT-Distill (red diamond) is on or above the Random/PAPIT-CLIP curve on 7 of 9 (dataset, retention) cells; at TextVQA k=25% ($\sim$10× FLOP reduction) it gains +12.3 over random and essentially recovers Unpruned. Annotations on each x-tick give the retention k.

Table 2. **Alignment diagnostic** between PAPIT-CLIP scores and LLaVA's last-layer LLM attention over 576 image tokens (100 samples per dataset). All three datasets give $\rho\approx0$ and top-10% Jaccard *below* the uniform-random reference of 0.053, so any CLIP-aligned scorer is at-best random downstream.

Metric	GQA	TextVQA	VQA v2
Spearman ρ mean	−0.018	−0.016	0.014
Spearman ρ median	0.008	0.019	0.010
% of samples $\rho>0.3$	0.0%	0.0%	0.0%
% of samples $\rho<0$	48%	47%	43%
Top-10% Jaccard	0.047	0.039	0.046

benchmark, and PAPIT-CLIP's top-10% patches overlap with the LLM's most-attended patches less than uniform-random selection would. This rules out an entire family of training-free pruners that score patches in any CLIP-aligned space, regardless of the specific scoring rule (we verified that cosine and value-feature variants are no better; see [10] for the dense-feature analogue).

5.3. Mid-layer LLM attention does carry the signal

We test the converse with a two-pass oracle: run LLaVA once on the full prompt, score by attention from non-image queries to image keys at LLM layer L, prune, and re-generate. Table 3 sweeps $L\in\{2,8,16,24,31\}$ on GQA. Mid-layer attention ($L\in[8,24]$) consistently beats random; layer 2 (FastV's recommended layer in its single-pass setting) and the last layer do not. This is the positive control that motivates distillation.

5.4. PAPIT-Distill: making the oracle deployable

The oracle costs an extra full LLaVA forward, defeating the efficiency goal. We instead distil it. **Architecture ablation** (Table 4) on the 5K Scope-A cache picks MLP_4 as the lightest predictor that beats both MLP_2 and a cross-attention block. **Supervision target ablation** on a 20K subset of the 100K cache (val-KL on a 10% held-out split, MLP_4, target $\in\{A^{(8)},A^{(16)},A^{(24)},\text{rollout}\}$) finds L=8 minimises val-KL (0.094) and maximises top-10% Jaccard; rollout is hardest to predict because it aggregates across

Table 3. **Two-pass LLM-attention oracle on GQA** (100 samples, exact match). **Bold**: beats random at that retention. Layer 16 gives the largest margin at low retention; FastV's layer 2 fails in this two-pass setup because the prompt-only forward has not yet committed attention mass to query-relevant patches.

Selector	k=25%	k=50%	k=75%
Unpruned	0.58	0.58	0.58
Random	0.53	0.53	0.56
PAPIT-CLIP	0.56	0.53	0.51
LLM-attn L=2	0.49	0.54	0.53
LLM-attn L=8	**0.58**	**0.54**	0.56
LLM-attn L=16	**0.61**	**0.58**	0.56
LLM-attn L=24	**0.60**	**0.57**	0.55
LLM-attn L=31	0.50	**0.55**	0.55

Table 4. **Predictor architecture ablation**, MLP_4 wins on Scope-A 5K held-out at supervision target L=16. All four architectures share the same KL training loop (Eq. 3).

Arch	Params	val-KL↓	top-10% Jac.↑
Bilinear	0.92M	0.43	0.34
MLP_2	0.59M	0.12	0.54
MLP_4	1.28M	**0.087**	**0.62**
Cross-attn	1.58M	0.089	0.61

all 32 layers [16]. Critically, downstream accuracy follows distillability rather than oracle strength: although Table 3 shows L=16 is the strongest oracle, the L=8 *predictor* wins the downstream ablation across the three benchmarks at N=100 (+0.030 over random, vs. +0.024 for L=16, +0.021 for L=24, and −0.054 for rollout — i.e. the rollout distill is *worse* than random, confirming that a hard-to-learn target produces a useless predictor regardless of its oracle quality). The deployed predictor is therefore MLP_4 trained against $A^{(8)}$.

Downstream accuracy of the deployed PAPIT-Distill is reported in the bottom rows of Table 1 and visualised in Figure 4. Per the latency benchmark in Section 3, the scorer adds 0.43 ms per image on top of the unmodified ViT forward, vs. 65 ms for PAPIT-CLIP (which incurs a

CLIP text-encoder forward pass and a backward pass for GradCAM). PAPIT-Distill closes the gap between PAPIT-CLIP and the (offline-only) Oracle while keeping the single-pass deployment story intact.

6. Discussion and Conclusion

Why doesn't PAPIT-CLIP collapse below random? LLaVA-1.5 exhibits substantial token redundancy [8], so at $k \geq 50\%$ the dominant accuracy driver is simply the *number* of retained tokens, not their identity. PAPIT-CLIP and Random therefore sit within ± 1 of each other on every cell. Layer 2 attention fails for a related reason: in our two-pass setup the prompt forward has not yet generated any answer token, so its attention is dominated by positional and register-like patterns [11]. Layers 8–24 have integrated the question with the image; layer 31 has already committed to generation rather than visual reading.

Limitations and outlook. Both diagnostics and downstream accuracy are reported at N=100 per (dataset, retention) cell; binomial $\sigma \approx 0.05$, so single-cell gaps under ± 0.10 are at most $\sim 2\sigma$. The TextVQA $k = 25\%$ gain (+12.3 over random) is therefore $\sim 2.5\sigma$ — large enough to take seriously but worth replicating at N=500+ before publication. Extending to LLaVA-1.6, Qwen2.5-VL, and single-pass FastV-style integration of PAPIT-Distill *inside* the LLM are natural next steps. The headline finding — a 1.3M-param MLP suffices to bridge a training-free, CLIP-derived scorer to one that respects what the LLM actually attends to — is the smallest viable fix to a clean negative result, and is the contribution we hope generalises beyond LLaVA-1.5.

Acknowledgements. We used GitHub Copilot for autocomplete and Claude (Anthropic) as a writing and code-review assistant. All experimental design, training runs, and analysis are our own.

References

[1] H. Liu, C. Li, Y. Li, and Y. J. Lee, "Improved baselines with visual instruction tuning," in *CVPR*, 2024.

[2] A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, S. Agarwal, G. Sastry, A. Askell, P. Mishkin, J. Clark, *et al.*, "Learning transferable visual models from natural language supervision," in *ICML*, 2021.

[3] A. Dosovitskiy, L. Beyer, A. Kolesnikov, D. Weissenborn, X. Zhai, T. Unterthiner, M. Dehghani, M. Minderer, G. Heigold, S. Gelly, J. Uszkoreit, and N. Houlsby, "An image is worth 16x16 words: Transformers for image recognition at scale," in *ICLR*, 2021.

[4] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin, "Attention is all you need," in *NeurIPS*, 2017.

[5] Y. Rao, W. Zhao, B. Liu, J. Lu, J. Zhou, and C.-J. Hsieh, "Dynamicvit: Efficient vision transformers with adaptive token sparsification," in *NeurIPS*, 2021.

[6] L. Meng, H. Li, B.-C. Chen, S. Lan, Z. Wu, Y.-G. Jiang, and S.-N. Lim, "Adavit: Adaptive vision transformers for efficient image recognition," in *CVPR*, 2022.

[7] D. Bolya, C.-Y. Fu, X. Dai, P. Zhang, C. Feichtenhofer, and J. Hoffman, "Token merging: Your ViT but faster," in *ICLR*, 2023.

[8] L. Chen, H. Zhao, T. Liu, S. Bai, J. Lin, C. Zhou, and B. Chang, "An image is worth 1/2 tokens after layer 2: Plug-and-play inference acceleration for large vision-language models," *arXiv preprint arXiv:2403.06764*, 2024.

[9] Y. Zhang, C.-K. Fan, J. Ma, W. Zheng, T. Huang, K. Cheng, D. Gudovskiy, T. Okuno, Y. Nakata, K. Keutzer, and S. Zhang, "SparseVLM: Visual token sparsification for efficient vision-language model inference," *arXiv preprint arXiv:2410.04417*, 2024.

[10] C. Zhou, C. C. Loy, and B. Dai, "Extract free dense labels from CLIP," in *ECCV*, 2022.

[11] T. Darcet, M. Oquab, J. Mairal, and P. Bojanowski, "Vision transformers need registers," in *ICLR*, 2024.

[12] R. R. Selvaraju, M. Cogswell, A. Das, R. Vedantam, D. Parikh, and D. Batra, "Grad-CAM: Visual explanations from deep networks via gradient-based localization," in *ICCV*, 2017.

[13] D. A. Hudson and C. D. Manning, "Gqa: A new dataset for real-world visual reasoning and compositional question answering," in *CVPR*, 2019.

[14] Y. Goyal, T. Khot, D. Summers-Stay, D. Batra, and D. Parikh, "Making the v in vqa matter: Elevating the role of image understanding in visual question answering," in *CVPR*, 2017.

[15] A. Singh, V. Natarajan, M. Shah, Y. Jiang, X. Chen, D. Batra, D. Parikh, and M. Rohrbach, "Towards vqa models that can read," in *CVPR*, 2019.

[16] S. Abnar and W. Zuidema, "Quantifying attention flow in transformers," in *ACL*, 2020.